

CS528

Caching and Multi-threading

A Sahu
Dept of CSE, IIT Guwahati

Program Cache Behavior: Hit/Miss

Cache model

- Direct mapped 8 word per line



Program

```
int A[128];  
for (i=0; i<128; i++) {  
    A[i]=i;  
}
```

- Assume &A=000000, **Behavior of only Data**
- Scalar variable {i} mapped to register
- Data have to moved from cache/memory

Cache perf. : Data Size <= Cache Size

```
int A[128];  
for (i=0; i<128; i++) {  
    A[i]=i;  
}
```

Scalar mapped to register
Vector mapped to memory

1:7= 1miss:7hit

1:7	0	A[0]	A[1]	A[2]					A[7]
2:14	1	A[8]	A[9]						A[15]
3:21	2	A[16]	A[17]						A[23]
	14								
16:112	15								A[127]

Strided access: Reduce locality

```
for (i=0; i<N; i++) {  
    for (j=0; j<N; j++) {  
        a[i][j]=i*j  
    }  
}    /* (a+i*N+j), j++
```

Row major
access: Stride 1,
improve locality,
cache hit

```
for (i=0; i<N; i++) {  
    for (j=0; j<N; j++) {  
        a[j][i]=i*j  
    }  
}    /* (a+j*N+i), j++
```

Column major
access: Stride N,
No locality, cache
miss dominates

Matrix mult.c

```
int A[8][8], B[8][8], C[8][8];
for (i=0; i<8; i++) {
    for (j=0; j<8; j++) {
        S=0;
        for (k=0; k<8; k++)
            S=S+B[i][k]*C[k][j];
        A[i][j]=S;
    }
}
```

Data Size > Cache Size

- $(64+64+64) > 128$ words
- When we get into cache it can take benefit
for $(k=0; k<8; k++)$

$S = S + \mathbf{B[i][k]} * \mathbf{C[k][j]};$

- Inner loop execute for **64 times**
 - We have to get $B[j]$ once will have 1miss/7 hit
 - $C[k]$ have to bring every time 8miss
 - Total = **7h+9m**
- 2nd loop A have one miss in 8 access (1miss/7hit)
 - Total for A = **8m+56h**
- Total program : $64*(7h+9m)+8m+56h=504h+584m$
- **Miss Probability = $584/(504+584)=0.5367$**

Improving Locality

Matrix Multiplication example

$$[C] = [A] \times [B]$$

$L \times M$

$L \times N$


$N \times M$

Cache Organization for the example

- Cache line (or block) = 8 matrix elements.
- Matrices are stored row wise.
- **Cache can't accommodate a full row/column.**
 - **L, M and N are so large w.r.t. the cache size**
 - After an iteration along any of the three indices, when an element is accessed again, it results in a miss.
- Ignore misses due to conflict between matrices.
 - As if there was a **separate cache for each matrix.**

Matrix Multiplication : Code I

```
for (i = 0; i < L; i++)  
  for (j = 0; j < M; j++)  
    for (k = 0; k < N; k++)  
      C[i][j] += A[i][k] * B[k][j];
```




	C	A	B
accesses	LM	LMN	LMN
misses	LM/8	LMN/8	LMN

Total misses = $LM(9N+1)/8$

$L=M=N=100$; miss = $100*100*901/8=1,126,250$

Matrix Multiplication : Code II

```
for (k = 0; k < N; k++)  
  for (i = 0; i < L; i++)  
    for (j = 0; j < M; j++)  
      C[i][j] += A[i][k] * B[k][j];
```




	C	A	B
accesses	LMN	LN	LMN
misses	LMN/8	LN	LMN/8

Total misses = $LN(2M+8)/8$

$L=M=N=100$; miss = $100*100*208/8=260,000$

Matrix Multiplication : Code III

```
for (i = 0; i < L; i++)  
  for (k = 0; k < N; k++)  
    for (j = 0; j < M; j++)  
      C[i][j] += A[i][k] * B[k][j];
```




	C	A	B
accesses	LMN	LN	LMN
misses	LMN/8	LN/8	LMN/8

Total misses = $LN(2M+1)/8$

$L=M=N=100$; miss = $100*100*201/8=251,250$

Matrix Multiplication : Code III

```
for (i = 0; i < L; i++)  
  for (k = 0; k < N; k++)  
    for (j = 0; j < M; j++)  
      C[i][j] += A[i][k] * B[k][j];
```



All most all modern processor uses

- Cache block pre-fetch
- When ith block is getting used i+1 block prefetched
- **Perfect overlap : only three cache miss**
 - **Each for A, B, C**



**See the cachegrind demo of
matmul code**

Case Study : Jacobi-Algorithm Stencil based iterative method

```
double D[2][kmax+1][kmax+1];
int t0=0,t1=1;
for(time=0; time<itertime; time++){
    for(i=1;i<kmax,i++){
        for(j=1;j<kmax;j++){
            //Sum of neighbours
            S= D[t0][i+1][j]+D[t0][i-1][j]
              + D[t0][i][j-1]+ D[t0][i][j+1];
            D[t1][i][j]=S*0.25;
        }
        tmp=t0; t0=t1; t1=tmp; //swap array
    }
}
```

Inner loop : $Bc = 5Words/4F = 1.25 W/F$

WriteNot Allocate: $Bc = 4W/4F = 1.0F$

Jacobi-Algorithm

```
double D[2][kmax+1][kmax+1];
int t0=0,t1=1;
for(time=0; time<itertime; time++){
    for(i=1;i<kmax,i++) {
        for(j=1;j<kmax;j++){
            //Sum of neighbours
            S= D[t0][i+1][j]+D[t0][i-1][j]
              + D[t0][i][j-1]+ D[t0][i][j+1];
            D[t1][i][j]=S*0.25;
        }
        tmp=t0; t0=t1; t1=tmp; //swap array
    }
}
```

**Assume Row i and i-1 with no cost : if cache
is capable to holding two rows
 $B_c = 1W/4F = 0.25W/F$**

Algorithm Classification and Access Optimization

- $O(N)/O(N)$: If the # of arithmetic Ops and data transfer (LD/ST) are proportional to Loop Length N
 - Optimization potential is limited
 - Example Scalar Product, vector add, sparse MVM
- Memory bound for large N
- Compiler generated code achieve good perf.
 - Using software pipelining and loop nests

Loop fusion for $O(N)/O(N)$

```
for (i=0; i<N; i++)  
    A[i]=B[i]+C[i];    //  $B_c=3W/1F$   
for (i=0; i<N; i++)  
    Z[i]=B[i]+E[i];    //  $B_c=3W/1F$ 
```



```
for (i=0; i<N; i++) {  
    A[i]=B[i]+C[i];  
    Z[i]=B[i]+E[i];  
}
```

$B_c=5W/2F$
No need to B[i]

$O(N^2)/O(N^2)$: OPS/DataTransfer

- Typical two loop nests with loop strip count N
 - $O(N^2)$ operation for $O(N^2)$ loads and stores
- Example: dense MVM, Mat add, MatTrans
- MVM : -> Covert both access to row access

```
for (i=0; i<N; i++) {  
    tmp=C[i]  
    for (j=0; j<N; j++)    tmp=A[i][j]*B[j]  
    C[i]=tmp  
}
```

- Row I of A and vector B
- Original $Bc=2W/2F$ but $\rightarrow 2W*m/2F$
- m is miss rate of cache for Row access

$O(N^3)/O(N^2)$: OPS/DataTransfer

- Typical three loop nests
 - $O(N^3)$ operation for $O(N^2)$ loads and stores
- Example: dense Matrix Multiplication
- Implementation of cache Bound
 - Already studied : loop interchange
 - **Blocking : Strassen multiplication, will be discussed later**

Threading

Threading Language and Support

- Pthread: POSIX thread
 - Popular, Initial and Basic one
- Improved Constructs for threading
 - c++ thread : available in c++11, c++14
 - Java thread : very good memory model
 - Atomic function, Mutex
- Thread Pooling and higher level management
 - OpenMP (loop based)
 - Cilk (dynamic DAG based)

Programming with Threads

- Threads
- Shared variables
- The need for synchronization
- Synchronizing with semaphores
- Thread safety and reentrancy
- Races and deadlocks

Traditional View of a Process

- Process = process context + code, data, and stack

Process context

Program context:

Data registers

Condition codes

Stack pointer (SP)

Program counter (PC)


Kernel context:

VM structures (VMem)

Descriptor table

brk pointer

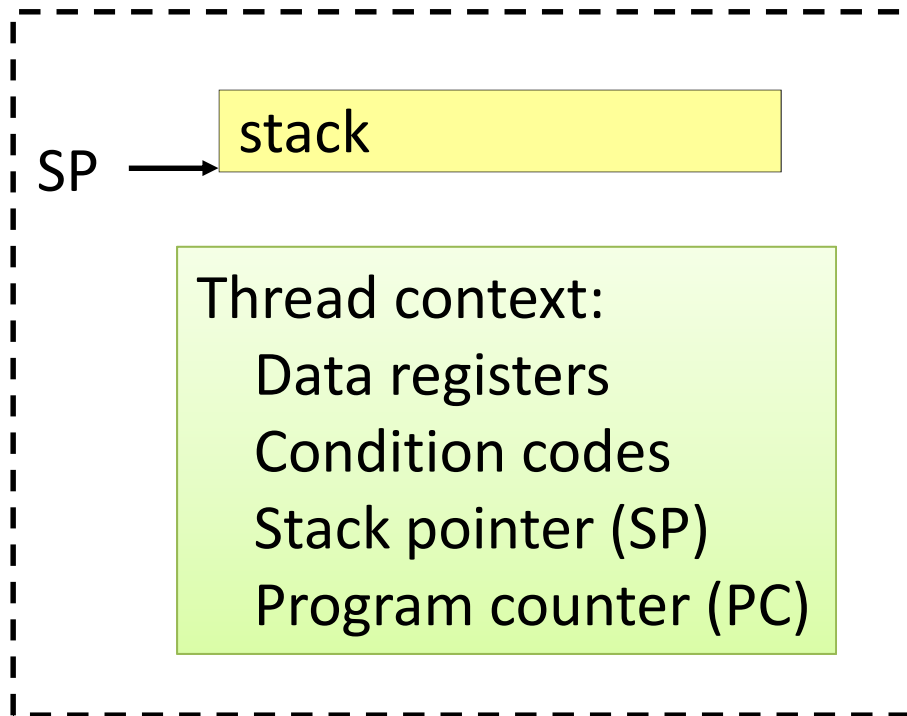
Code, data, and stack



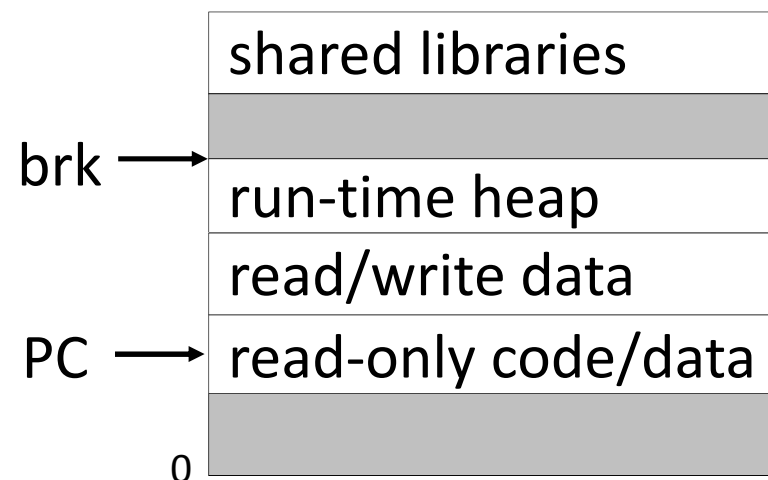
Alternate View of a Process

- Process = thread+ code, data & kernel context

Thread (main thread)



Code and Data

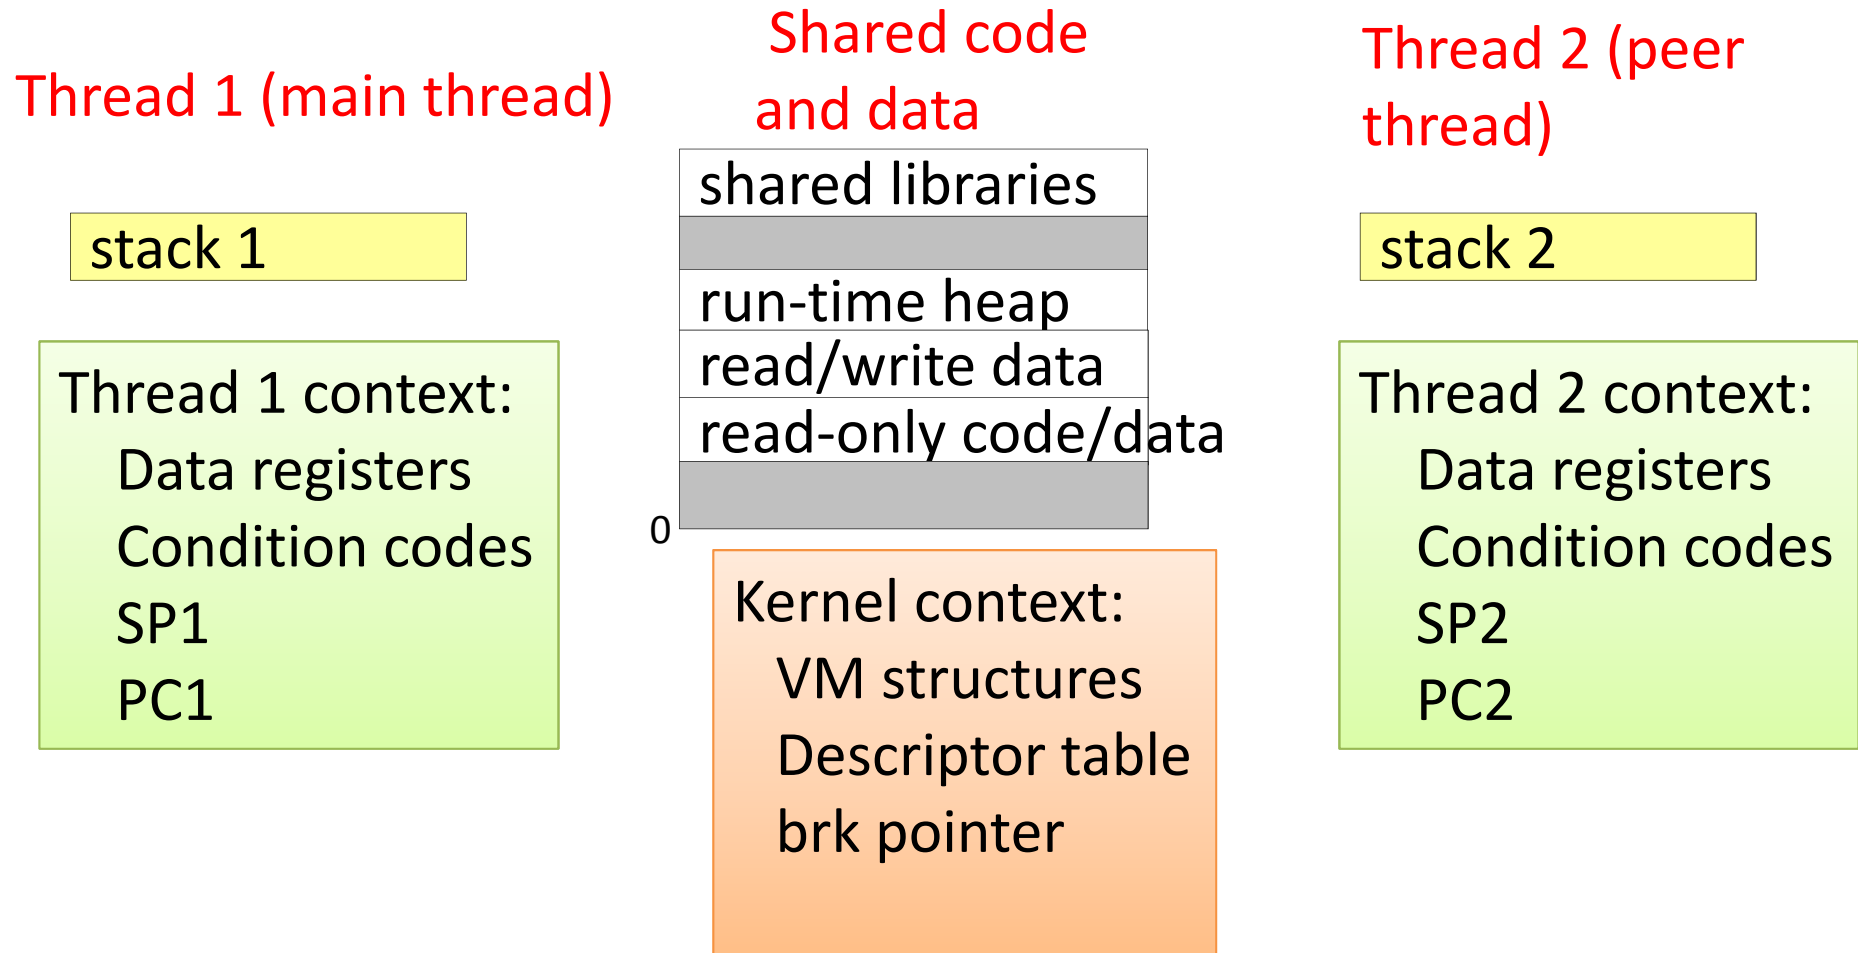


Kernel context:
VM structures
Descriptor table
brk pointer

A Process With Multiple Threads

- Multiple threads can be associated with a process
 - Each thread has its *own* logical control flow (sequence of PC values)
 - Each thread *shares* the same code, data, and kernel context
 - Each thread has its own thread id (TID)


A Process With Multiple Threads



Logical View of Threads

- Threads associated with a process form a pool of peers
 - Unlike processes, which form a tree hierarchy

Threads associated with
process foo



Process hierarchy

