

CS528

Multi-threading

A Sahu
Dept of CSE, IIT Guwahati

Threading Language and Support

- Pthread: POSIX thread
 - Popular, Initial and Basic one
- Improved Constructs for threading
 - c++ thread : available in c++11, c++14
 - Java thread : very good memory model
 - Atomic function, Mutex
- Thread Pooling and higher level management
 - OpenMP (loop based)
 - Cilk (dynamic DAG based)

Pthread, C++ Thread, Cilk and OpenMP

```
pthread_t tid1, tid2;
pthread_create(&tid1, NULL, Fun1, NULL);
pthread_create(&tid2, NULL, Fun2, NULL);
pthread_join(tid1, NULL);
pthread_join(tid2, NULL);
```

Pthread

```
thread t1(Fun1);
thread t2(Fun2, 0, 1, 2); // 0, 1, 2 param to Fun2
t1.join();
t2.join();
```

C++
thread

```
#pragma omp parallel for
```

```
for(i=0; i<N; i++)
```

```
    A[i]=B[i]*C[i];
```

```
//Auto convert serial code to threaded code
```

```
// $gcc -fopenmp test.c; export OMP_NUM_THREADS=10; ./a.out
```

```
cilk fib (int n){//Cilk dynamic parallism, DAG recursive code
    if (n<2) return n;
    int x=spawn fib(n-1); //spawn new thread
    tnt y=spawn fib(n-2); //spawn new thread
    sync;
    return x+y;
}
```

Programming with Threads

- Threads
- Shared variables
- The need for synchronization
- Synchronizing with semaphores
- Thread safety and reentrancy
- Races and deadlocks

Traditional View of a Process

- Process = process context + code, data, and stack

Process context

Program context:

Data registers

Condition codes

Stack pointer (SP)

Program counter (PC)


Kernel context:

VM structures (VMem)

Descriptor table

brk pointer

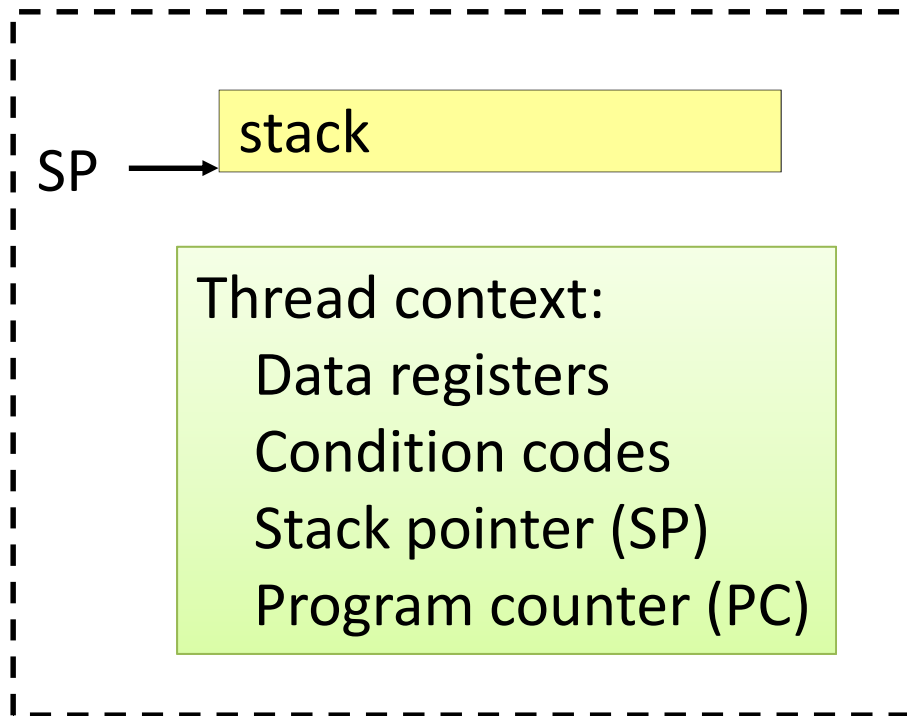
Code, data, and stack




Alternate View of a Process

- Process = thread+ code, data & kernel context

Thread (main thread)



Code and Data

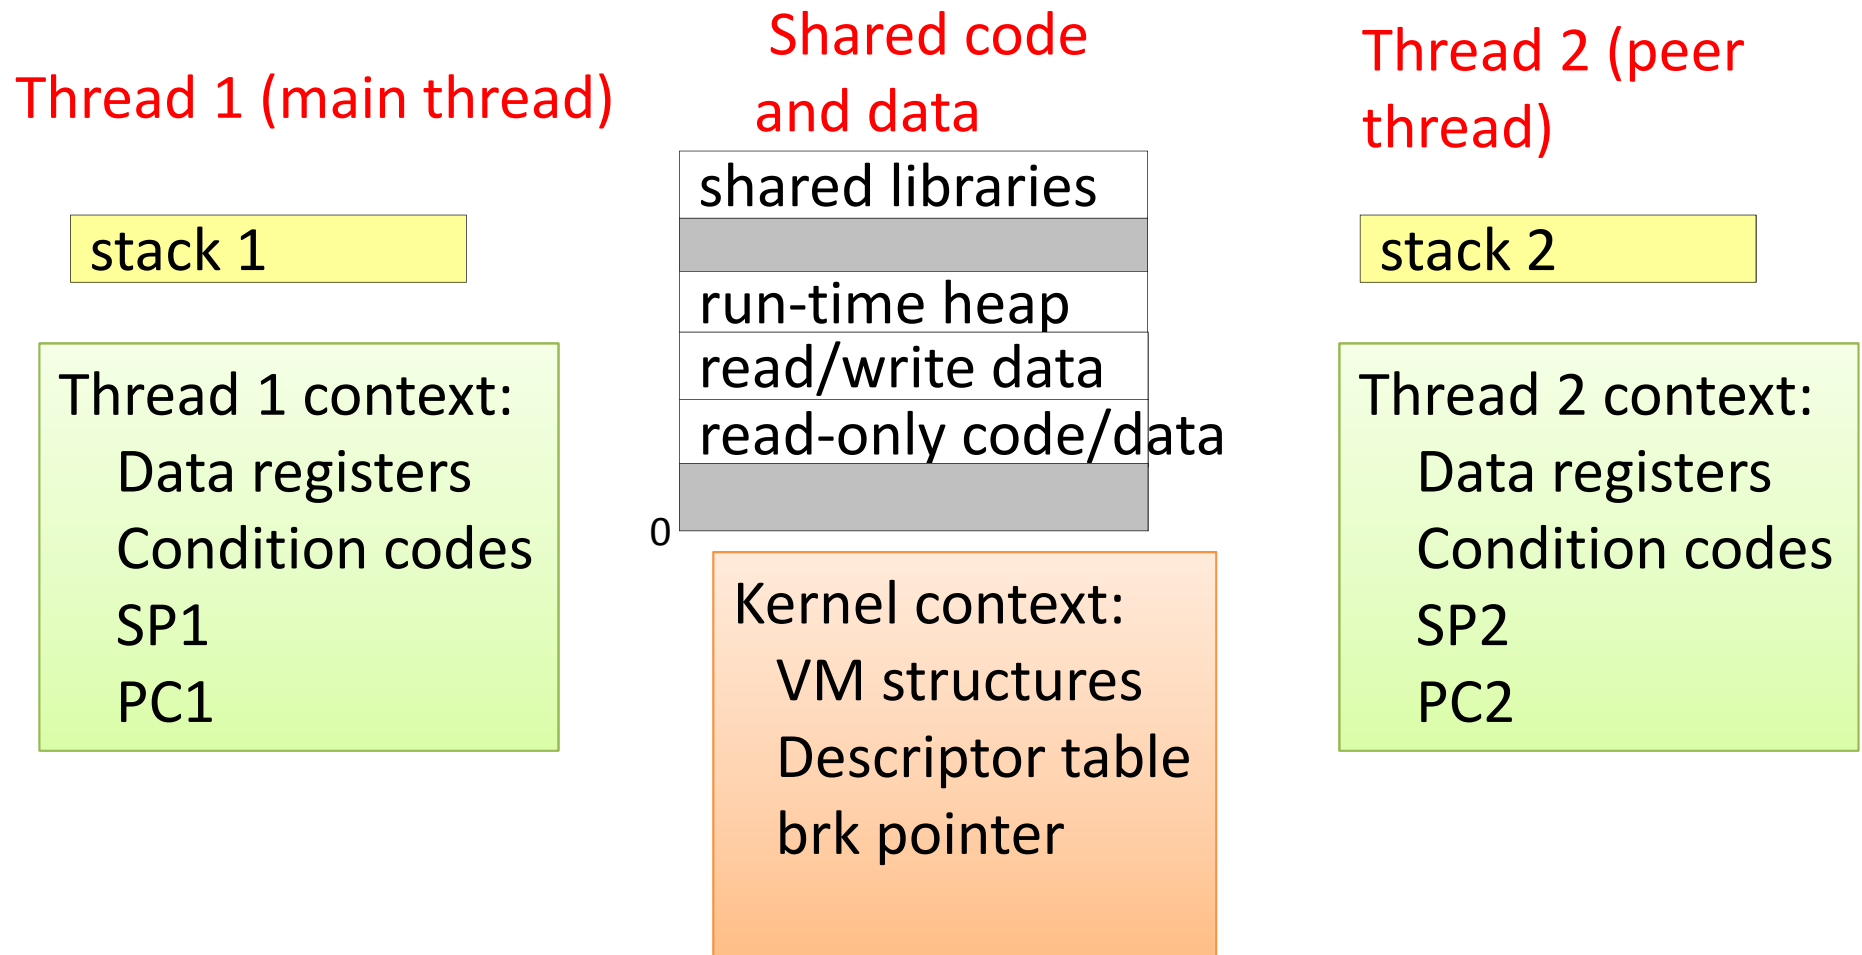


Kernel context:
VM structures
Descriptor table
brk pointer

A Process With Multiple Threads

- Multiple threads can be associated with a process
 - Each thread has its *own* logical control flow (sequence of PC values)
 - Each thread *shares* the same code, data, and kernel context
 - Each thread has its own thread id (TID)


A Process With Multiple Threads




Logical View of Threads

- Threads associated with a process form a pool of peers
 - Unlike processes, which form a tree hierarchy

Threads associated with
process foo



Process hierarchy



Threads vs. Processes

- How threads and processes are similar
 - Each has its own logical control flow
 - Each can run concurrently
 - Each is context switched

Threads vs. Processes

- How threads and processes are different
 - Threads share code and data, processes (typically) do not
 - Threads are somewhat less expensive than processes
 - Process control (creating and reaping) is twice as expensive as thread control
 - Linux/Pentium III numbers:
 - ~20K cycles to create and reap a process
 - ~10K cycles to create and reap a thread

Posix Threads (Pthreads) Interface

- **Creating and reaping threads**
 - `pthread_create`, `pthread_join`
- **Determining your thread ID** : `pthread_self`
- **Terminating threads**
 - `pthread_cancel`, `pthread_exit`
 - `exit` [terminates all threads], `return` [terminates current thread]
- **Synchronizing access to shared variables**
 - `pthread_mutex_init`,
`pthread_mutex_[un]lock`
 - `pthread_cond_init`,
`pthread_cond_[timed]wait`

The Pthreads "hello, world" Program

```
/* thread routine */  
void *HelloW(void *vargp) {  
    printf("Hello, world!\n");  
    return NULL;  
}
```

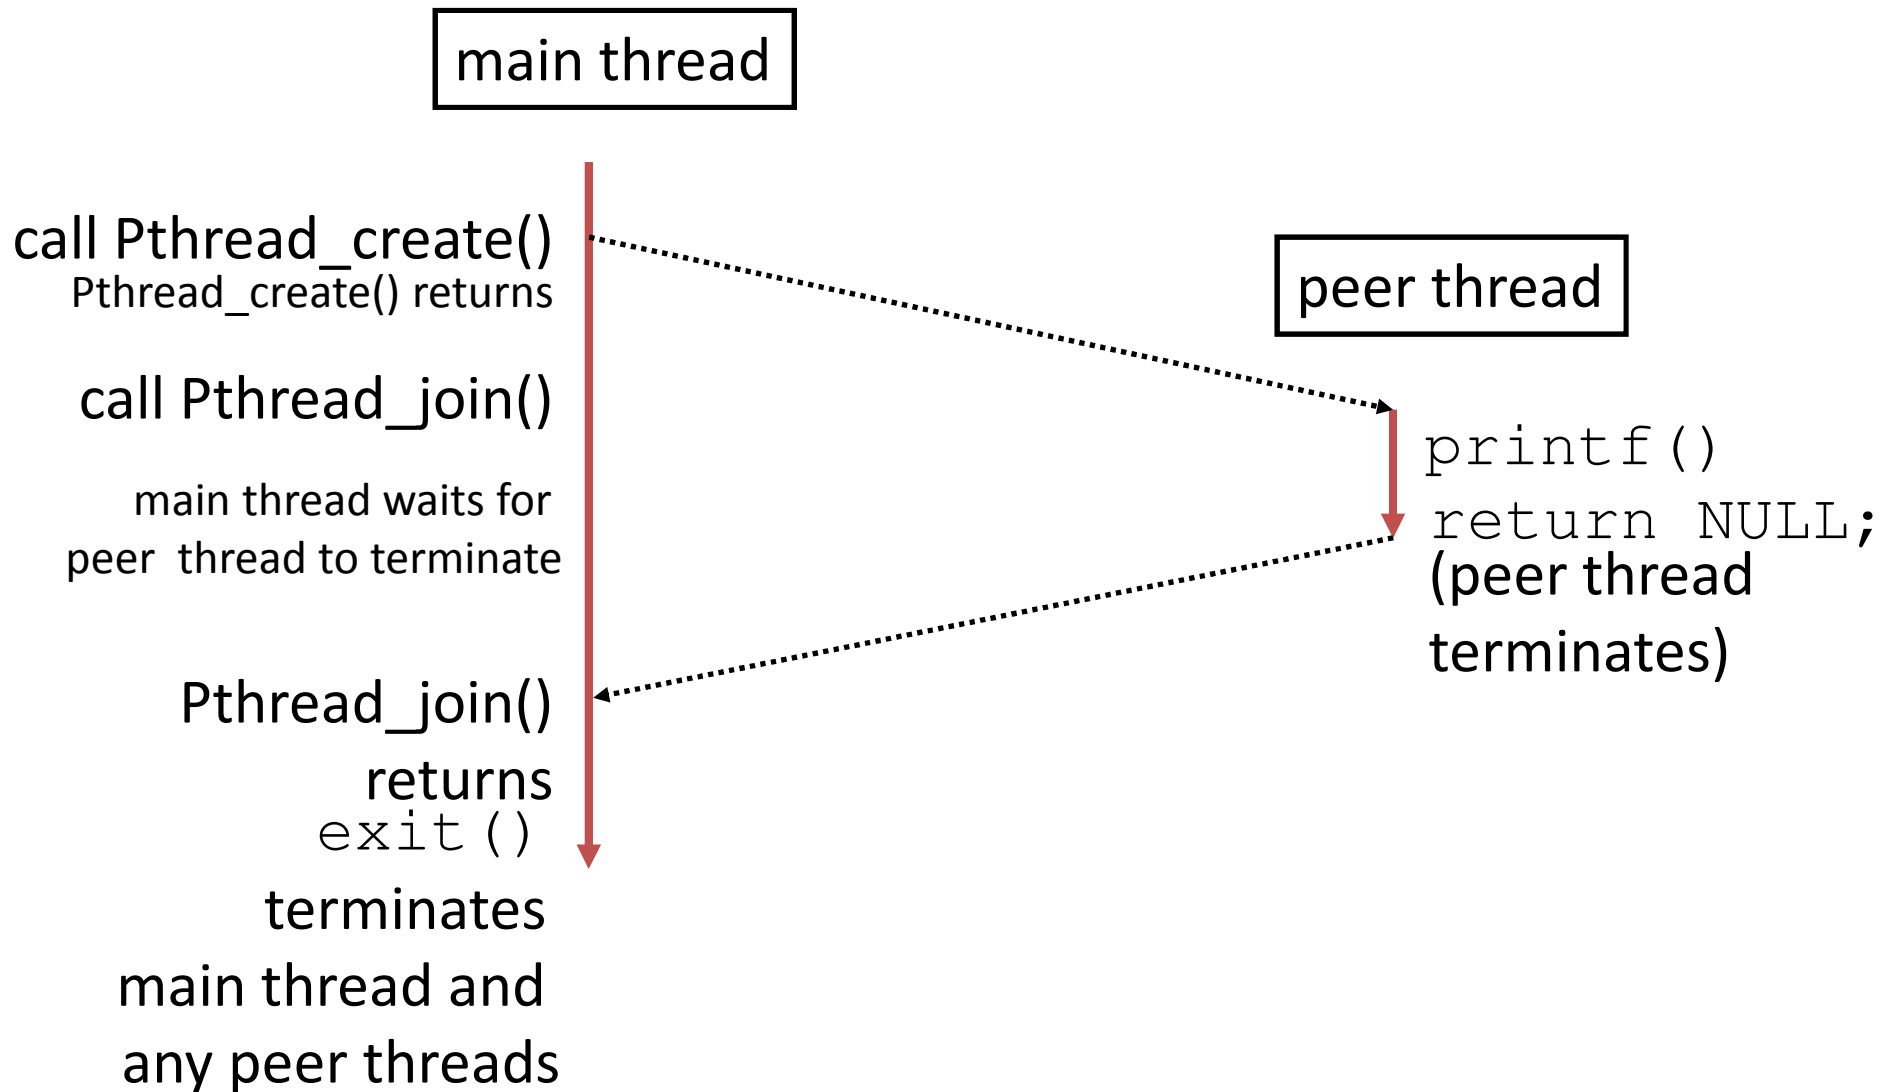
```
int main() {  
    pthread_t tid;  
    pthread_create(&tid, NULL, HelloW, NULL);  
    pthread_join(tid, NULL);  
    return 0;  
}
```

Thread attributes
(usually NULL)

Thread arguments
(void *p)

return value
(void **p)

Execution of Threaded “hello, world”



Pros and Cons: Thread-Based Designs

- **+ Easy to share data structures between threads**
 - E.g., logging information, file cache
- **+ Threads are more efficient than processes**
- **– Unintentional sharing can introduce subtle and hard-to-reproduce errors!**
 - Ease of data sharing is greatest strength of threads
 - Also greatest weakness!

VectorSum Serial

```
int A[VSize], B[VSize], C[VSize];  
  
void VectorSumSerial() {  
    for( int j=0; j<SIZE; j++)  
        A[j]=B[j]+C[j];  
}
```

Suppose Size=1000

0-249	250-499	500-749	750-999
--------------	----------------	----------------	----------------

T1

T2

T3

T4

VectorSum Serial

```
int A[VSize], B[VSize], C[VSize];

void VectorSumSerial() {
    for( int j=0; j<SIZE; j++)
        A[j]=B[j]+C[j];
}
```

- Independent
- Divide work into equal for each thread
- Work per thread: $\text{Size}/\text{numThread}$

VectorSum Parallel

```
void *DoVectorSum(void *tid) {  
    int j, SzPerthrd, LB, UB, TID;  
    TID= *((int *)tid);  
    SzPerthrd= (VSize/NUM_THREADS);  
    LB= SzPerthrd*TID; UB=LB+SzPerthrd;  
  
    for (j=LB; j<UB; j++)  
        A[j]=B[j]+C[j];  
}
```



VectorSum Parallel

```
int main() {  
    int i;  
    pthread_t thread[NUM_THREADS];  
    for (i = 0; i < NUM_THREADS; i++)  
        pthread_create(&thread[i],  
            NULL, DoVectorSum, (void*)&i);  
    for (i = 0; i < NUM_THREADS; i++)  
        pthread_join(thread[i], NULL);  
  
    return 0;  
}
```

Matrix multiply and threaded matrix multiply

- Matrix multiply: $C = A \times B$

$$C[i, j] = \sum_{k=1}^N A[i, k] \times B[k, j]$$



Matrix multiply and threaded matrix multiply

- Matrix multiply: $C = A \times B$

$$C[i, j] = \sum_{k=1}^N A[i, k] \times B[k, j]$$

- Divide the whole rows to T chunks
 - Each chunk contains : N/T rows, Assume $N\%T=0$

Matrix multiply Serial

```
void MatMul () {  
    int i, j, k, S;  
    for (i=0; i<Size; i++)  
        for (j=0; j<Size; j++) {  
            S=0;  
            for (k=0; k<Size; k++)  
                S=S+A[i][k]*B[k][j];  
            C[i][j]=S;  
        }  
}
```

Matrix Pthreaded: RowWise

```
void * DoMatMulThread(void *arg) {  
    int i, j, k, S, LB, UB, TID, ThrdSz;  
    TID = * ((int *) arg); ThrdSz = Size / NumThrd;  
    LB = TID * ThrdSz; UB = LB + ThrdSz;  
  
    for (i = LB; i < UB; i++)  
        for (j = 0; j < Size; j++) {  
            S = 0;  
            for (k = 0; k < Size; k++)  
                S = S + A[i][k] * B[k][j];  
            C[i][j] = S;  
        }  
}
```

Matrix Pthreaded: RowWise

```
int main() {
    pthread_t  thread[NumThread];
    int t;
    Initialize();
    for (t=0; t<NumThread; t++)
        pthread_create(&thread[t], NULL,
            DoMatMulThread, &t);
    for (t=0; t<NumThread; t++)
        pthread_join(thread[t], NULL);
    TestResult();
    return 0;
}
```


Estimating π using Monte Carlo


- The probability of a random point lying inside the unit circle:

$$P(x^2 + y^2 < 1) = \frac{A_{\text{circle}}}{A_{\text{square}}} = \frac{\pi}{4}$$

- If pick a random point N times and M of those times the point lies inside the unit circle:

$$P^{\circ}(x^2 + y^2 < 1) = \frac{M}{N}$$

- If N becomes very large, $P = P^{\circ}$



$$\pi = \frac{4 \cdot M}{N}$$

Value of PI: Monte-Carlo Method

```
void MontePI () {  
    int count=0,i;  
    double x,y,z;  
    for ( i=0; i<niter; i++) {  
        x = (double)rand() /RAND_MAX;  
        y = (double)rand() /RAND_MAX;  
        z = x*x+y*y;  
        if (z<=1) count++;  
    }  
    pi=(double) count/niter*4;  
}
```

PI- Multi-threaded

- 1 thread you are able to generate N points
 - Suppose M points fall under unit circle
 - $PI = 4M/N$
- With 10 thread generate 10XN points and calculate more accurately
 - Each thread calculate own value of PI (or M)
 - Average later on (or recalculate PI from collective M)

Value of PI: Pthreaded

```
int main() {  
    pthread_t  thread[NumThread]; double pi;  
    int t, at[NumThread], count, TotalIter;  
    for(t=0; t<NumThread; t++)  
        pthread_create(&thread[t], NULL,  
            DoLocalMC_PI, &t);  
    for(t=0; t<NumThread; t++)  
        pthread_join(thread[t], NULL);  
    for(t=0; t<NumThread; t++) count+=LCount[t];  
    TotalIter=niter*NumThread;  
    pi=((double) count/TotalIter) *4;  
    return 0;  
}
```

Value of PI: Pthreaded

```
int LCount [NumThread];  
void *DoLocalMC_Pi (void *aTid) {  
    int tid, count, i; double x, y, z;  
    tid= *((int *) aTid);  
    count=0; LCount[tid]=0;  
    for ( i=0; i<niter; i++) {  
        x = (double) rand() / RAND_MAX;  
        y = (double) rand() / RAND_MAX;  
        z = x*x+y*y; if (z<=1) count++;  
    }  
    LCount[tid]=count;  
}
```

What is the sieve of Eratosthenes?

- Used to find prime number between 2 and N
- It works by gradually eliminating multiple of smallest unmark prime (x) in the given interval
 - Till $x^2 > N$

Sieve of Eratosthenes

- Current Prime 2

[2,3,4,5,6,7,8,9,10,11,12,13, 14,15,16,17,18,19,20]

- Current Prime 2

[2,3,4,5, ~~6~~,7,8,9,~~10~~,11,~~12~~,13,~~14~~,15,~~16~~,17,~~18~~,19,~~20~~]

- Current Prime 3

List : [2, 3, 5, 7, 9, 11, 13, 15, 17, 19]

- Current Prime 3

List : [2, 3, 5, 7, ~~9~~, 11,13,~~15~~, 17,19]

- Current Prime 5 and $5 = \text{ceil}(\text{sqrt}(20))$

List : [2, 3, 5, 7, 11, 13, 17,19] //All are primes

Sieve of Eratosthenes

```
int prime[MaxNum], CurrPrime;
void DoStriking() {
    for(int i=2*CurrPrime; i<MaxNum;
        i=i+CurrPrime)    prime[i]=0;
}
int main() {
    int i, sqrtmaxnum=sqrt(MaxNum)+1;
    for(i=0; i<MaxNum; i++)    prime[i]=1;
    for(i=2; i<sqrtmaxnum; i++) {
        if (prime[i]==0)    continue ;
        CurrentPrime=i; DoStriking();
    }
}
```


Sieve of Eratosthenes: Pthreaded

```
void *DoStriking(void *tid) {
    int i, SizePerThrd, LB, UB;
    int TID= *((int *)tid);
    SizePerThrd = (MaxNum/NUM_THREADS);
    LB= SizePerThrd * TID;
    UB=LB + SizePerThrd;
    if (LB< (2*CurrPrime) ) LB=2*CurrPrime;
    for (i=LB; i<UB; i=i+CurrPrime)
        prime[i]=0;
}
```

Sieve of Eratosthenes: Pthreaded

```
for (i=2; i<sqrtmaxnum; i++) {  
    if (prime[i]==0) continue ;  
    CurrPrime=i;  
    for (j = 0; j < NUM_THREADS; j++)  
        pthread_create(&thread[j],  
            NULL, DoStriking, (void*)&j);  
    for (j = 0; j < NUM_THREADS; j++)  
        pthread_join(thread[j], NULL);  
}
```

Shared Variables in Threaded C Programs

- Question: Which variables in a threaded C program are shared variables?
 - Answer not as simple as “global variables are shared” and “stack variables are private”
- Requires answers to the following questions:
 - What is the memory model for threads?
 - How are variables mapped to memory instances?
 - How many threads reference each of these instances?