

CS528

Multi-threading and OpenMP

A Sahu
Dept of CSE, IIT Guwahati

Outline

- Thread Memory Model
- Thread Safety : 4 Classic cases
- Thread Synchronization: Example
- Thread Programming Model
 - Boss/Worker, Peers, Pipeline
- Thread Pooling: Example : Manual
- **Implicit/Auto Thread Pooling: OpenMP/Cilk**

Thread Memory Model & Safety

Shared Variables in Threaded C Programs

- Question: Which variables in a threaded C program are shared variables?
 - Answer not as simple as “global variables are shared” and “stack variables are private”
- Requires answers to the following questions:
 - What is the memory model for threads?
 - How are variables mapped to memory instances?

Threads Memory Model

- Conceptual model:
 - Each thread runs in the context of a process
 - Each thread has its own separate thread context
 - Thread ID, stack, stack pointer, program counter, condition codes, and general purpose registers
 - All threads share remaining process context
 - Code, data, heap, and shared library segments of process virtual address space
 - Open files and installed handlers

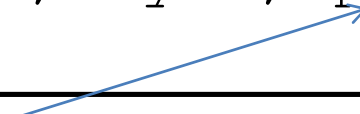
Threads Memory Model

- Operationally, this model is not strictly enforced:
 - Register values are truly separate and protected
 - But any thread can read and write the stack of any other thread
- *Mismatch between conceptual and operational model is a source of confusion and errors*

Threads Accessing other Thrd's Stack

```
char **ptr; /* global */
int main() {
    int i; pthread_t tid;
    char *msgs[N] = { "Hello1", "Hello2" };
    ptr = msgs;
    for (i = 0; i < 2; i++)
        pthread_create(&tid, NULL, thread, (void *)i);
    pthread_exit(NULL);
}
```

```
void *thread(void *vargp) {
    int myid = (int)vargp;
    printf("[%d]: %s \n", myid, ptr[myid]);
}
```



*Peer threads access main thread's stack
indirectly through global ptr variable*

Shared Variable Analysis

Variable instance	Ref by main thread?	Ref by peer thread 0?	Ref by peer thread 1?
<code>ptr</code>	yes	yes	yes
<code>i.m</code>	yes	no	no
<code>msgs.m</code>	yes	yes	yes
<code>myid.p0</code>	no	yes	no
<code>myid.p1</code>	no	no	yes

Answer: A variable `x` is shared iff multiple threads reference at least one instance of `x`. Thus:

- `ptr` and `msgs` are shared.
- `i` and `myid` are **NOT** shared.

Parallel Counter: without Lock

```
#define NITERS 100
int cnt = 0; /* shared */
int main() {
    pthread_t tid1, tid2;
    pthread_create(&tid1, NULL, count, NULL);
    pthread_create(&tid2, NULL, count, NULL);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    if (cnt != NITERS*2) printf("BOOM! cnt=%d", cnt);
    else printf("OK cnt=%d\n", cnt);
}
```

```
void *count(void *arg) {
for (int i=0; i<NITERS; i++) cnt++;
}
```

cnt should be 200
What went wrong?!

```
$/badcnt
BOOM! cnt=196
$/badcnt
BOOM! cnt=184
```

Thread Safety

- Functions called from a thread must be *thread-safe*
- There are four (non-disjoint) classes of thread-unsafe functions:
 - **Class 1: Failing to protect shared variables : L/UL**
 - Class 2: Relying on persistent state across invocations
 - Class 3: Returning pointer to static variable
 - Class 4: Calling thread-unsafe functions

Class 1: Failing to protect shared variables

- Fix: Use Lock and unlock semaphore operations
- Issue: Synchronization operations will slow down code
- Example: `goodcnt.c`

```
void *count(void *arg) {  
  for(int i=0; i<NITERS; i++)  
    pthread_mutex_lock(&LV);  
    cnt++;  
    pthread_mutex_unlock(&LV);  
} // LV is lock variable
```

Class 2: Relying on persistent state across multiple function invocations

- Random number generator relies on static state
- Fix: Rewrite function so that caller passes in all necessary state, → Maintain Thread Specific State

```
int rand() {  
    static uint next = 1;  
    next = next*1103515245 + 12345;  
    return (uint) (next/65536)% 32768;  
}  
void srand(uint seed) {  
    next = seed;  
}
```

Class 3: Returning pointer to static variable

- Fixes: 1. Rewrite code so caller passes pointer to `struct`, Issue: Requires changes in caller and callee
- *Lock-and-copy*: Issue: Requires only simple changes in caller (and none in callee), However, caller must free memory

```
struct hostent *gethostbyname (
                                char* name) {
    static struct hostent h;
    <contact DNS and fill in h>
    return &h;
}
```

Class 3: Returning pointer to static variable

```
struct hostent *gethostbyname_ts(char *p) {  
    struct hostent *q = Malloc(...);  
    P(&mutex); /* lock */  
    p = gethostbyname(name);  
    *q = *p;    /* copy */  
    V(&mutex);  
    return q;  
}
```

```
hostp = malloc(...);  
gethostbyname_r(name, hostp);
```

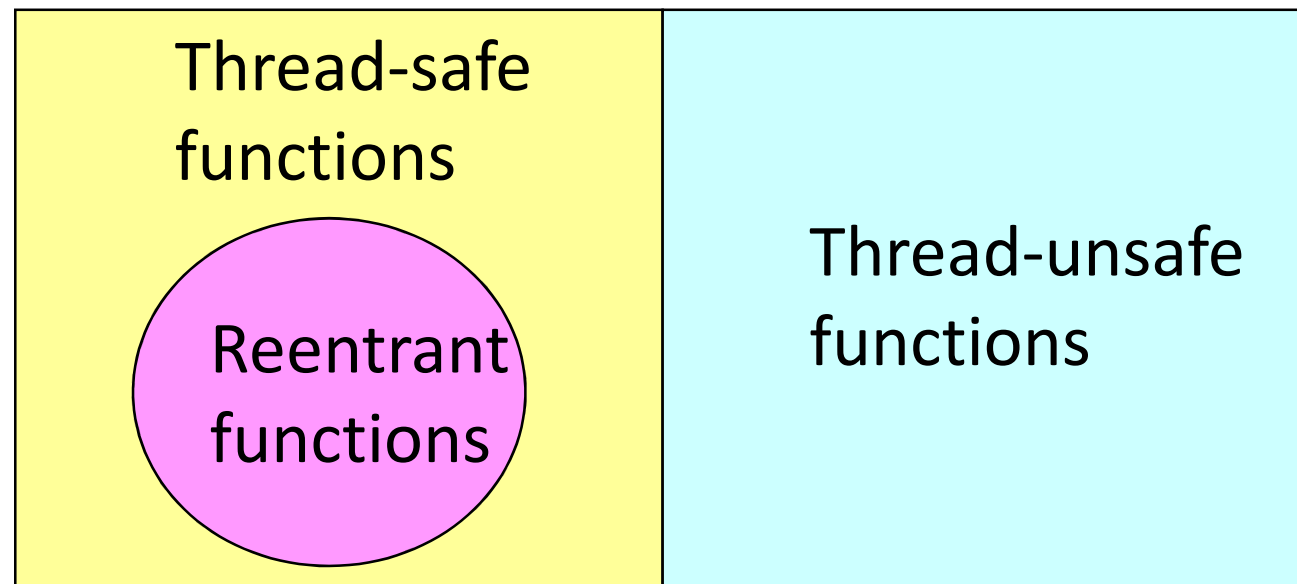
Class 4: Calling thread-unsafe functions

- Calling one thread-unsafe function makes an entire function thread-unsafe
- Fix: Modify the function so it calls only thread-safe functions

Reentrant Functions

- A function is *reentrant* iff it accesses NO shared variables when called from multiple threads
 - Reentrant functions are a proper subset of the set of thread-safe functions
 - NOTE: The fixes to Class 2 and 3 thread-unsafe functions require modifying the function to make it reentrant (only first fix for Class 3 is reentrant)

All functions



Thread-Safe Library Functions

- Most functions in the Standard C Library (at the back of your K&R text) are thread-safe
 - Examples: **malloc**, **free**, **printf**, **scanf**
- All Unix system calls are thread-safe
- Library calls that aren't thread-safe:

Thread-unsafe function	Class	Reentrant version
asctime	3	asctime_r
ctime	3	ctime_r
gethostbyaddr	3	gethostbyaddr_r
gethostbyname	3	gethostbyname_r
inet_ntoa	3	(none)
localtime	3	localtime_r
rand	2	rand_r

Synchronization : Lock/Unlock

Synchronization Hardware

- Modern machines provide special atomic hardware instructions
 - **Atomic** = non-interruptible
 - Either test memory word and set value
 - Or swap contents of two memory words
- Lock variable can be cached multicore
 - Cache coherence provide the correct Global state of LV
 - LL/SC

Atomic Sync Instructions

- Test and Set (TAS)
- Compare and Swap (CAS)
- Exchange (XCHG)
- Fetch and Increment (FAI)
- How to provide these
 - Load Locked and Store Conditional

test_and_set Instruction : TAS

```
//Definition:  
bool test_and_set (bool *LVptr) {  
    bool rv = *LVptr;  
    *LVptr = TRUE;  
    return rv;  
}
```

- Executed atomically
- Returns the original value of passed parameter
- Set the new value of passed parameter to “TRUE”.

CS : test_and_set Instruction : TAS

- Shared Boolean variable lock, initialized to FALSE
- Solution:

```
do {  
    while (test_and_set (&lock))  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = false;  
    /* remainder section */  
} while (true);
```

CAS: Compare and Set

```
int CAS(int *value, int expected, int
        new_value) {
    int temp = *value;
    if (*value == expected)
        *value = new_value;
    return temp;
}
```

- Executed atomically
- Returns the original value of passed parameter “value”
- Set the variable “value” the value of the passed parameter “new_value” but only if “value” == “expected”. That is, the swap takes place only under this condition.

Sync Solution using CAS


- Shared integer “lock” initialized to 0;

```
do {  
    while (CAS(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
    lock = 0;  
    /* remainder section */  
  
} while (true);
```


Synchronization Hierarchy 😊 😊 😊

- One == (used by) == > other
- **LL+SC ==> TAS/CAS/FAI/XCHG==>Lock/Unlock**
 - All **TAS/CAS/GAS/FAI/XCHG** do the same work
- Lock/Unlock == > Mutex //Mutex use L/UL
- Mutex == > Semaphore // Semaphore uses Mutex
 - Wait() and Signal()
- Semaphore == > Monitor //Monitor uses Semaphore
 - Many wait/Many Signal, Processes in Queue
 - Monitor : Another Abstract Type
 - *which use semaphore, mutex, conditions*

Many threads trying to acquire Mutex LOCK



Synchronization Primitives


- **int** pthread_mutex_init(
pthread_mutex_t *mutex_lock,
const pthread_mutexattr_t *lock_attr);
- **int** pthread_mutex_lock(
pthread_mutex_t *mutex_lock);
- **int** pthread_mutex_unlock(
pthread_mutex_t *mutex_lock);
- **int** pthread_mutex_trylock(
pthread_mutex_t *mutex_lock);

Locking Overhead

- Serialization points
 - Minimize the size of critical sections
 - Be careful
- Rather than wait, check if lock is available
 - **pthread_mutex_trylock**
 - If already locked, will return EBUSY
 - Will require restructuring of code
 - **Suspend self by pthread_yield() Give chance to others**
 - **Suspend self by doing a timed wait..**
 - {1, 1, 1,...}, {1, 2, 3, 4,...}, {1,2,4,8,...},...

Performance of Locking

- **Spinning/busy wait waste time**
- Recall MAC Protocol
 - Non Persistence CSMA protocol
 - Wait random time if medium if busy, then send
- Spin lock with exponential back-off reduces contention
 - Wait k amount of time for 1st attempt
 - Wait $k * c^i$ amount of time for i^{th} attempt



Example: simple backup

```
void ExpBackup(int _ival) {  
    int i=0, ival=_ival;  
    while (i<100) {  
        if (m.try_lock()) { ival=_ival;  
            BigNonSharedWork(); i++; m.unlock();  
        } else { ival=ival*2;  
            chrono::milliseconds interval(ival);  
            this_thread::sleep_for(interval);  
        } //End Else  
    } //End while  
} //EndExpBackup
```

Example :simple backup

```
int main(int argc, char *argv[]) {  
    thread th[10];  
    for(int i=0;i<10;i++)  
        th[i]=thread(ExpBackup,  
                     atoi(argv[1]));  
    for(int i=0;i<10;i++)  
        th[i].join();  
    return 0;  
}
```

Synch. Primitives

- `pthread_mutex_init, lock, unlock, trylock`
- `pthread_attr_setdetachstate, guardsize_np, stacksize, inheritsched, schedpolicy, schedparam`
- `pthread_cond_wait, signal, broadcast, init, destroy`
- Our main concern
 - Lock, unlock, trylock, condsignal, condbroadcast

Condition Variables for Synchronization

- Condition variable allows a thread
 - To block itself until specified data reaches a predefined state.
- Condition variable
 - Associated with a predicate (P)
 - When the predicate becomes true, the condition variable is used to signal one or more threads waiting on the condition.
- Single condition variable
 - May be associated with more than one predicate
 - Ex: $P = X \text{ OR } Y \text{ AND } (Z \text{ OR } K)$

C++ Thread:atomic

```
atomic_uint AtomicCount;
```

```
void DoCount () {
```

```
    int j, timesperthrd;
```

```
    timesperthrd= (TIMES/NUM_THREADS) ;
```

```
    for (j=0; j<timesperthrd; j++) AtomicCount++;
```

```
}
```

```
main () {
```

```
    thread T[N_THRDS]; int i;
```

```
    for (i=0; i<N_THRDS; i++) T[i]=thread(DoCount) ;
```

```
    for (i=0; i<N_THRDS; i++) T[i].join() ;
```

```
}
```

Improved

```
atomic_uint AtomicCount;
```

```
void DoCount () {
```

```
    int j, timesperthrd, localcount=0;
```

```
    timesperthrd=(TIMES/NUM_THREADS);
```

```
    for (j=0; j<timesperthrd; j++) localcount++;
```

```
        AtomicCount+=localcount;
```

```
}
```

```
main () {
```

```
    thread T[N_THRDS]; int i;
```

```
    for (i=0; i<N_THRDS; i++) T[i]=thread(DoCount);
```

```
    for (i=0; i<N_THRDS; i++) T[i].join();
```

```
}
```